

# Dargent: A Silver Bullet for Verified Data Layout Refinement (POPL 2023)

Chen-Lafont-O'Connor-Keller-McLaughlin-Jackson-Rizkallah

Dresden, 2<sup>nd</sup> June 2023

# Context

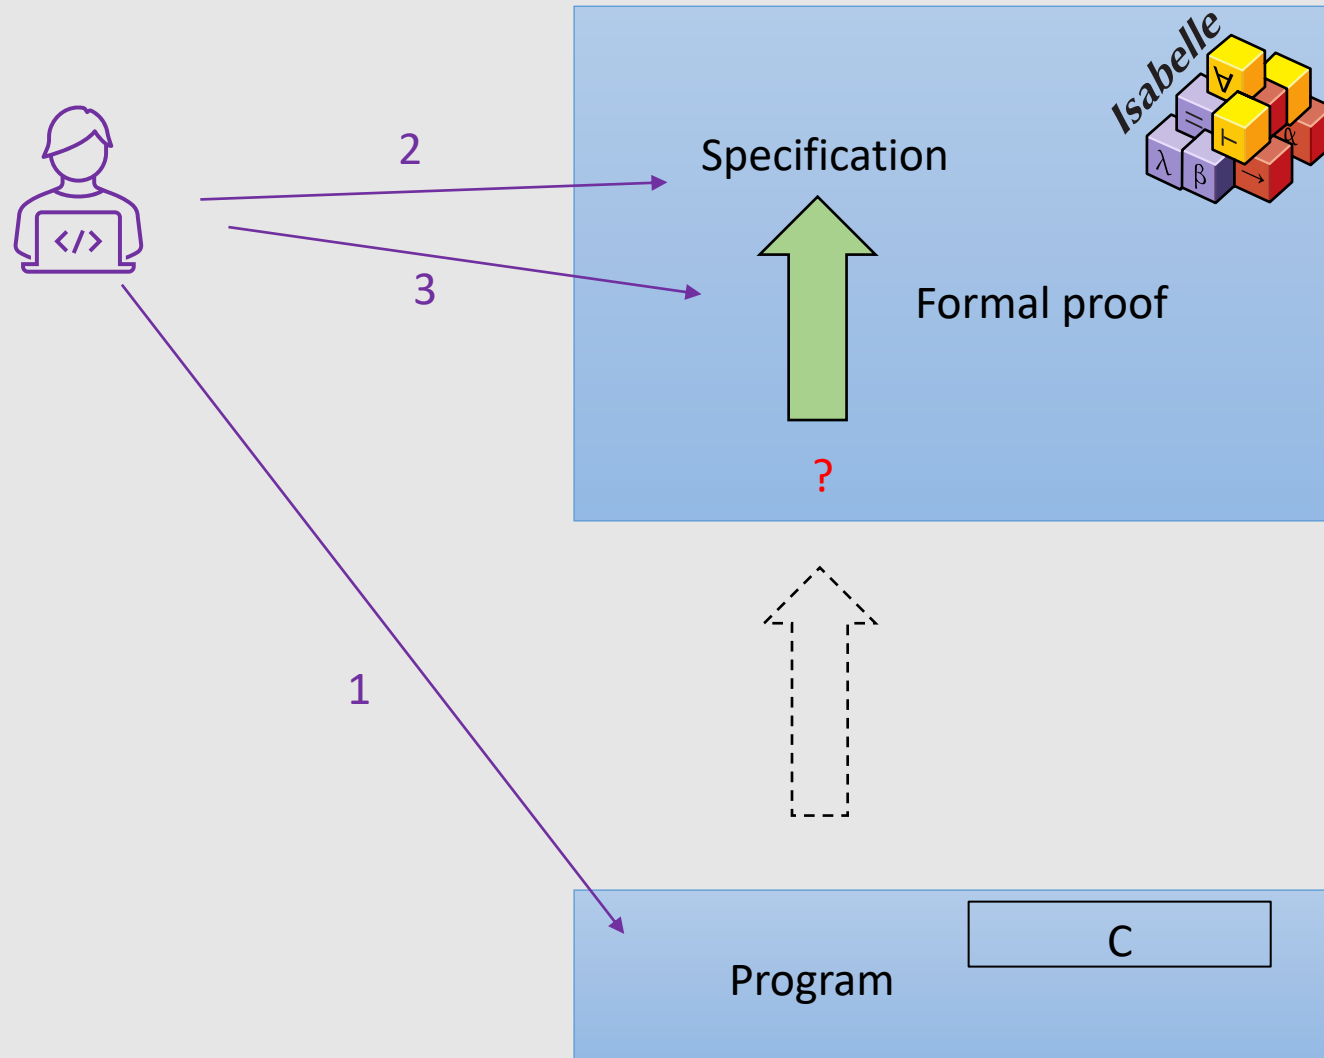


The world's most secure operating system kernel.

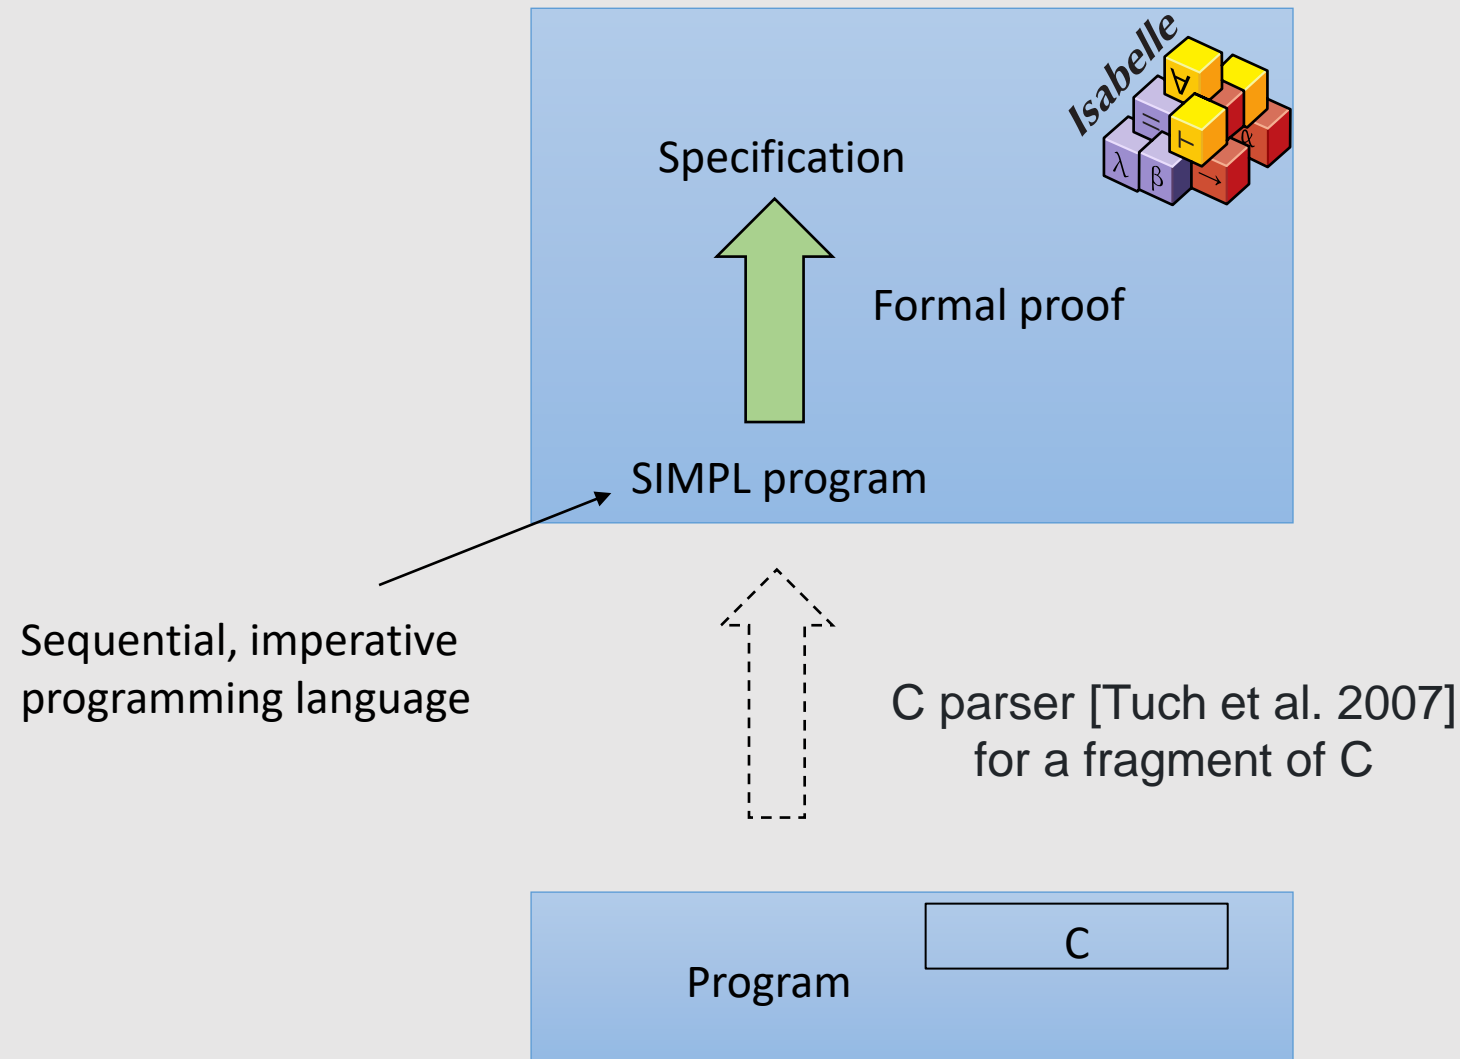


Used in  
Autonomous vehicles

# How does the sel4 team reason on C code?





# How does the sel4 team reason on C code?





Micro-kernel: 10k lines of C

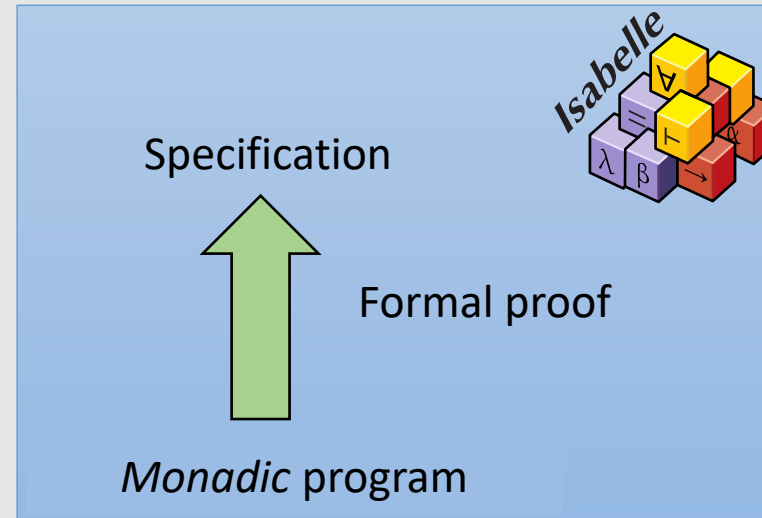


# 30M lines of C


File systems, device drivers, ...

How to ease systems code  
certification?

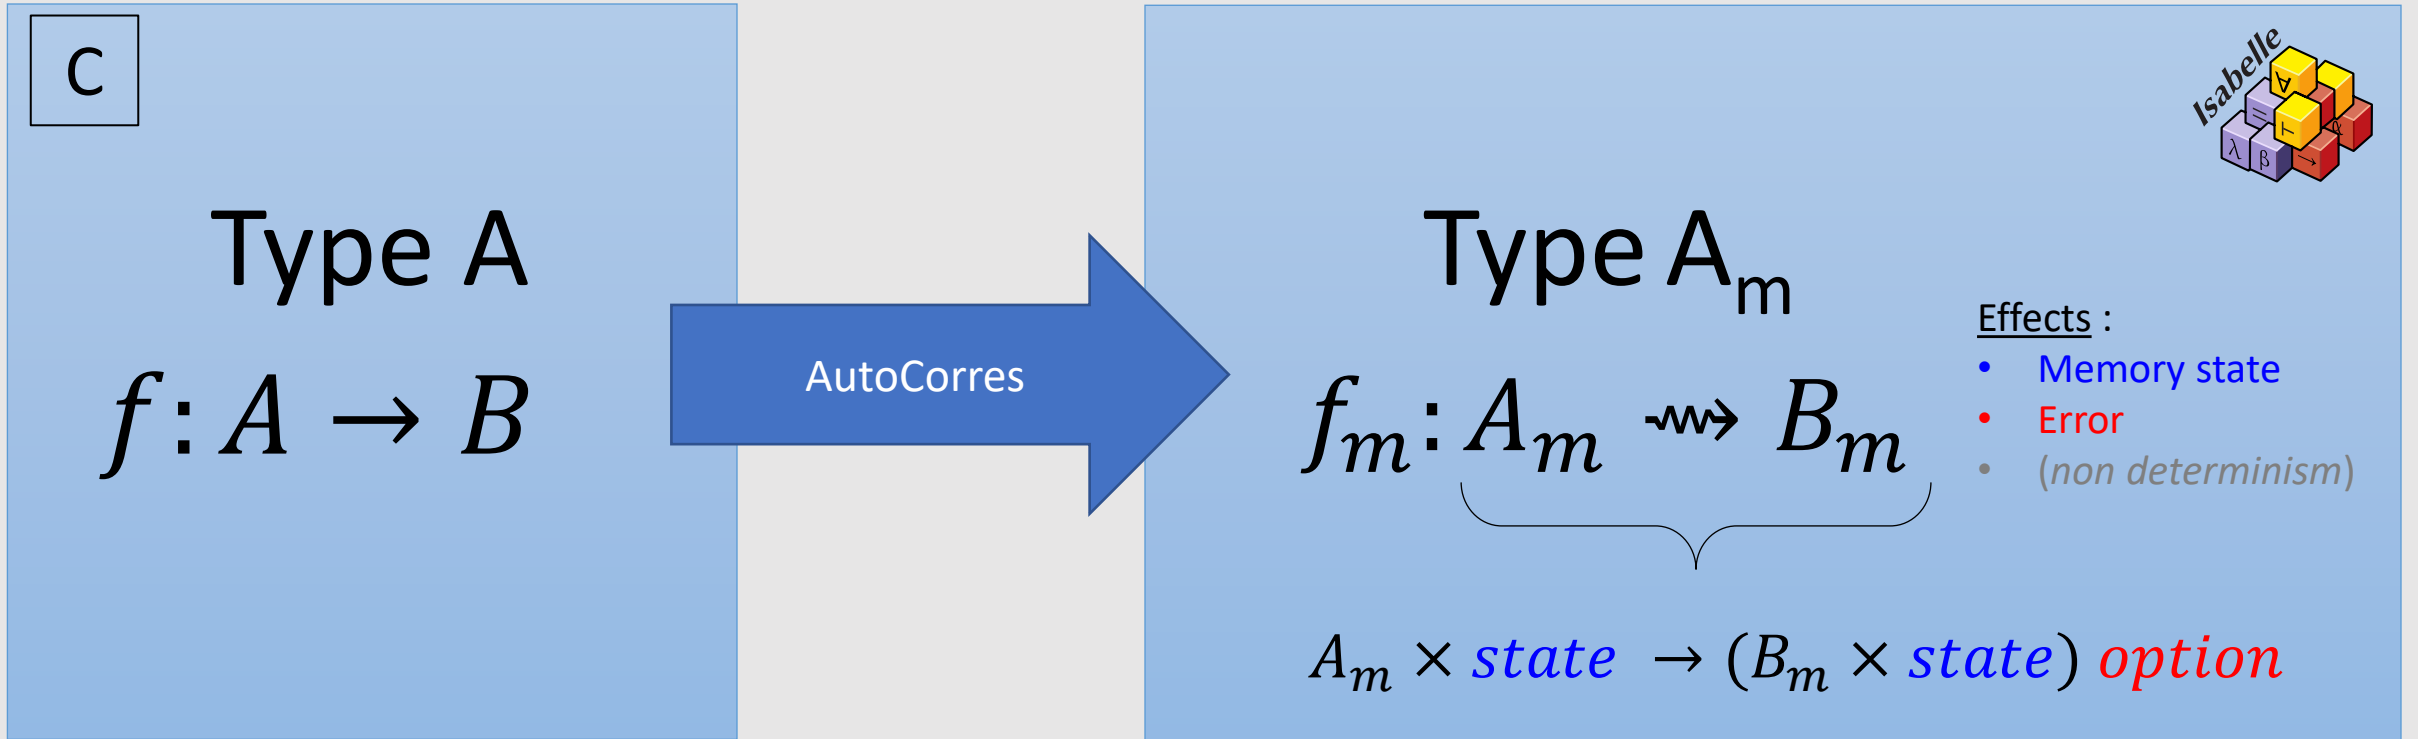
# A simpler embedding of C



**AutoCorres** [Greenaway et al. 2012]  
computes *monadic* semantics  
for a subfragment of the C parser

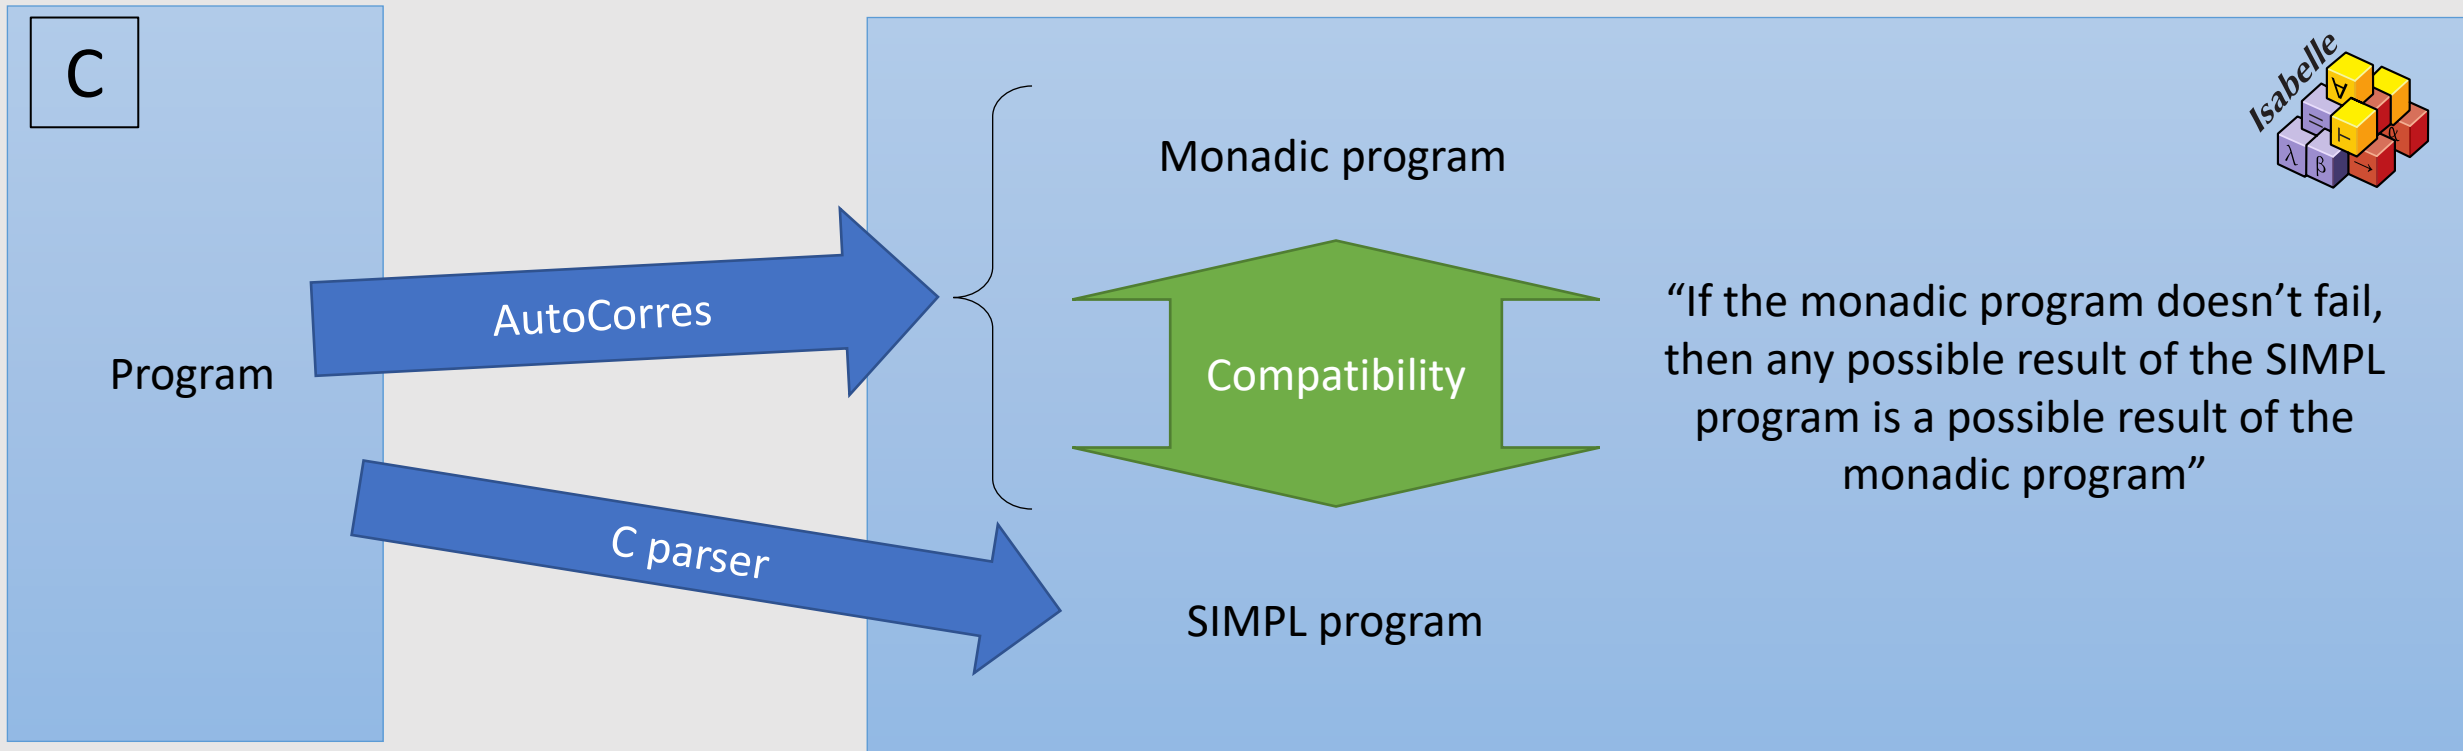


# The monadic semantics of AutoCorres





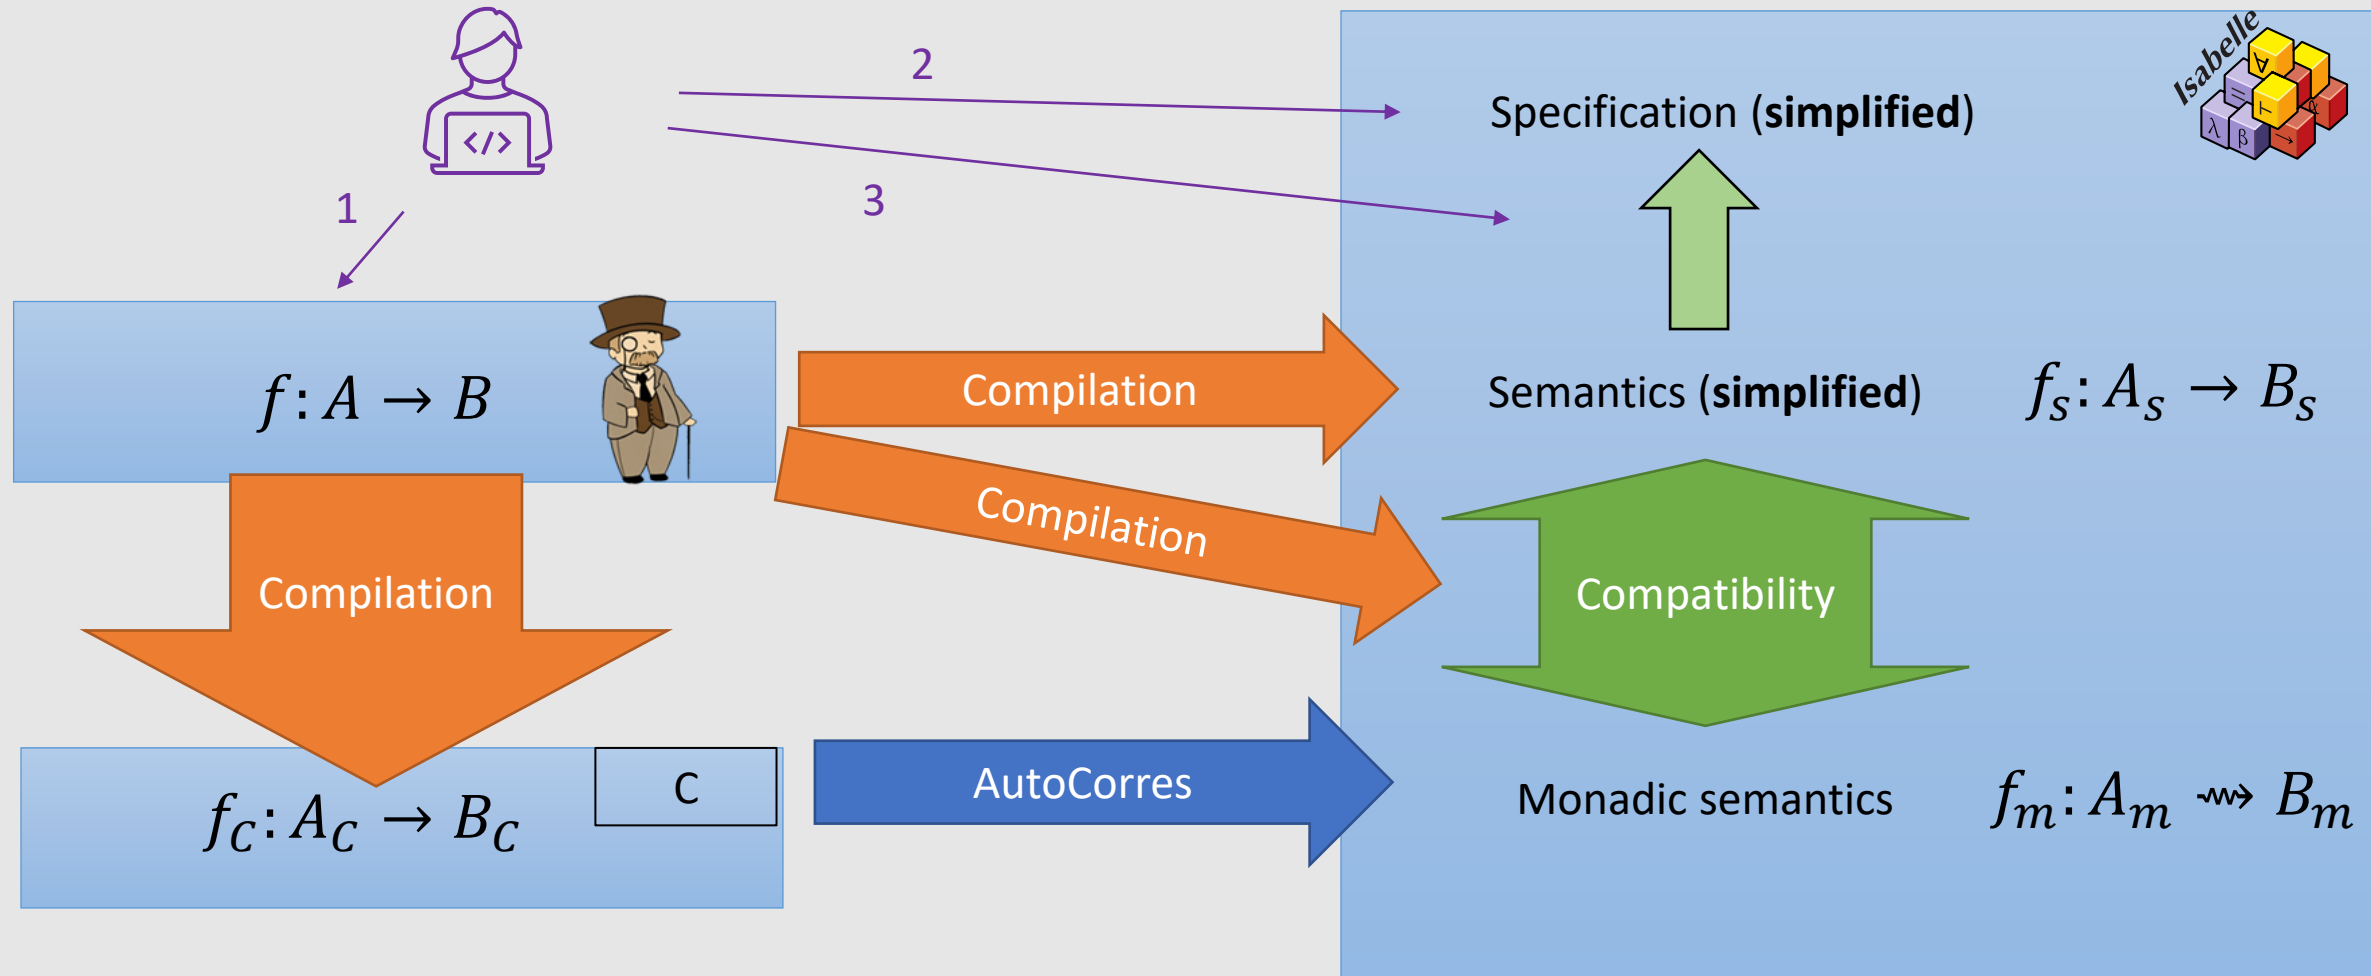
# Justifying the monadic semantics





# Cogent: Simplified semantics

A total functional language to describe a **safe** subfragment of AutoCorres, with **simplified semantics**



# Comparing semantics

Monadic semantics	Simplified semantics
Type $T_m$	Type $T_s$
$f_m: A_m \rightsquigarrow B_m$	$f_s: A_s \rightarrow B_s$

## Simplifications :

- No pointer / memory

$T_m = A_m^*$	$T_s = A_s$
---------------	-------------

- Pure functional semantics

$f_m: int \rightsquigarrow int$	$f_s: int \rightarrow int$
---------------------------------	----------------------------

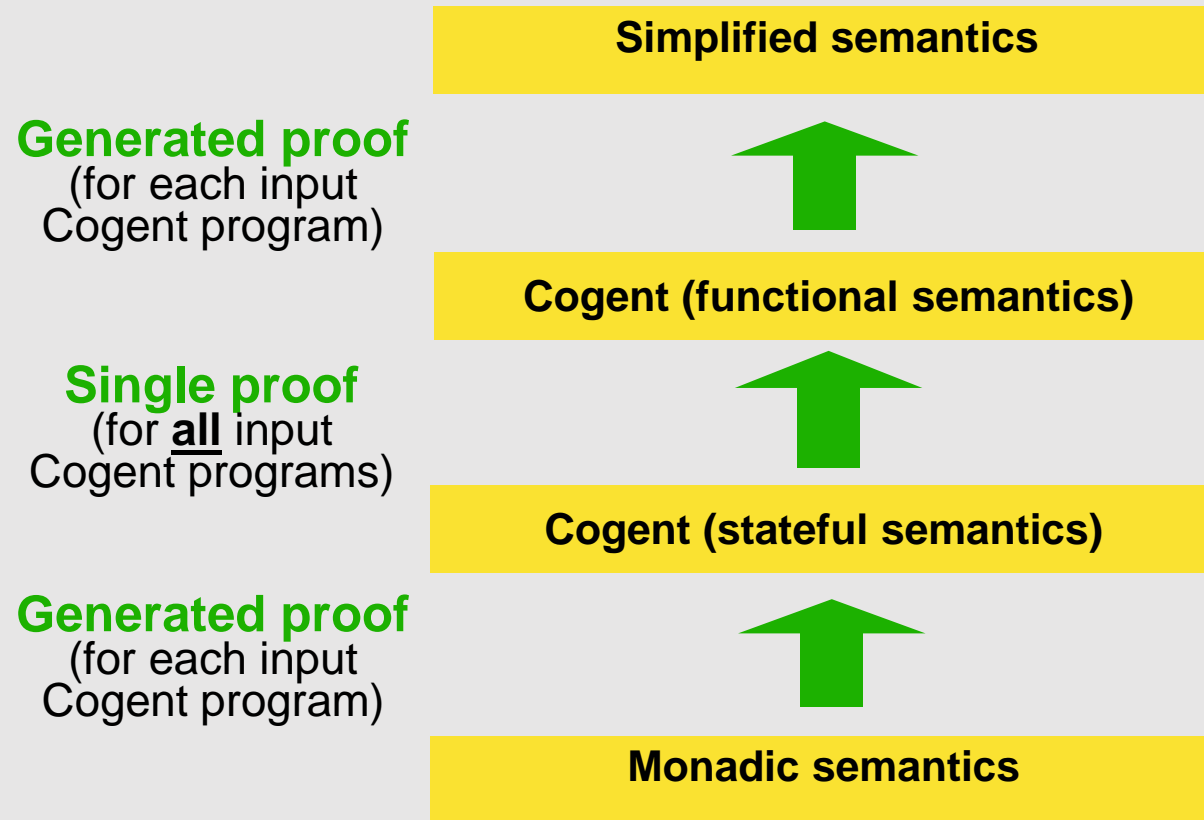
Compatibility between the two semantics
Relation $T_r \subset state \times T_m \times T_s$
$(f_m, f_s)$ is <i>compatible</i> with $(A_r, B_r)$

«  $(f_m, f_s)$  maps related values to related values »  
(in particular,  $f_m$  doesn't fail)

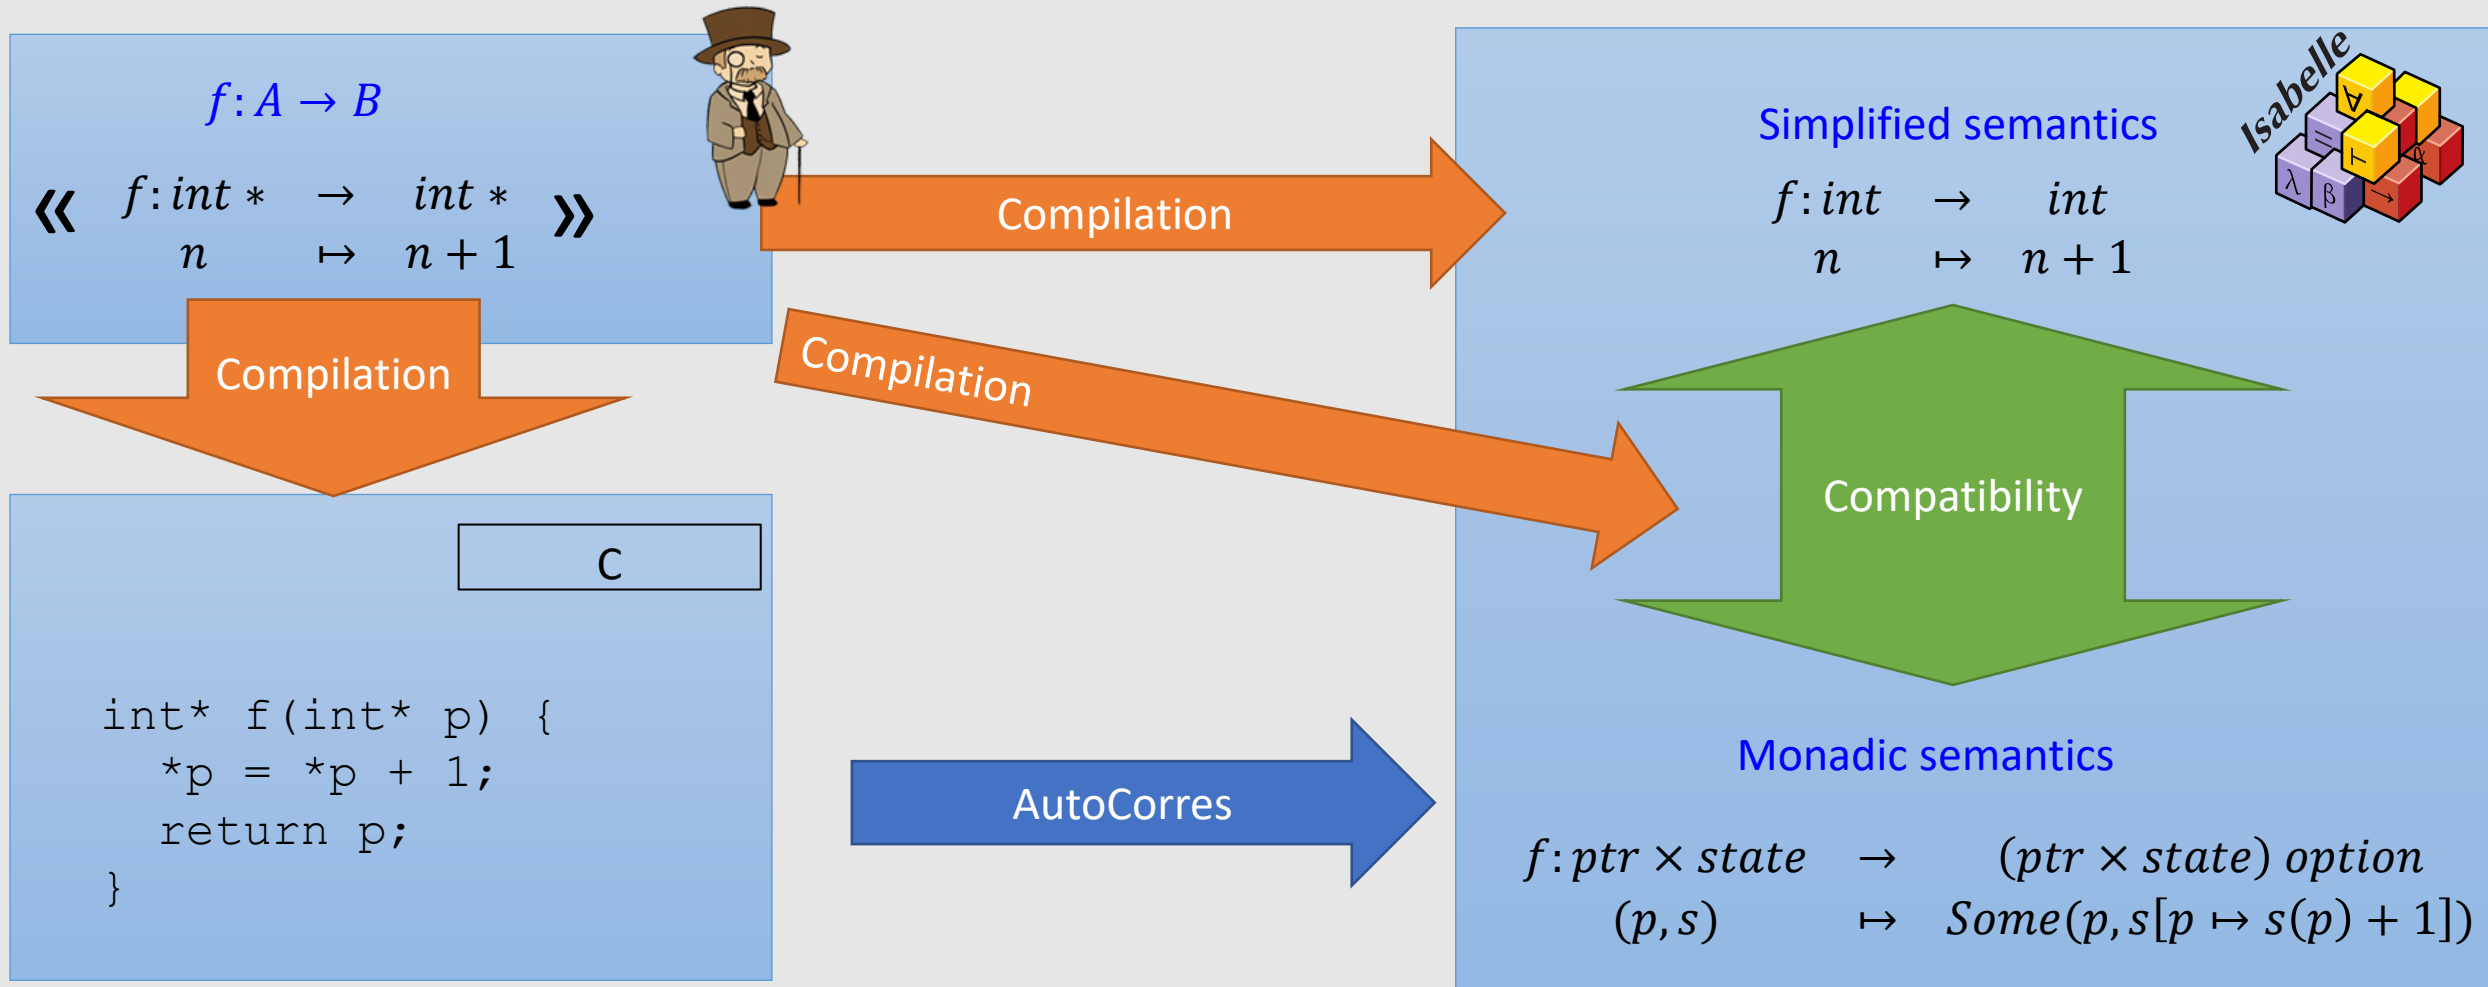
$T_r = \{(memory, p, memory(p))\}$
------------------------------------

$f_m(s, n) = Some(s', f_s(n))$
--------------------------------

# Compatibility, under the hood



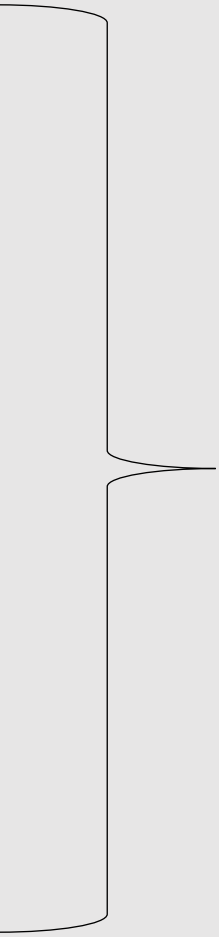
# Example



# Cogent

A total functional programming language with FFI support

- ▶ words (unsigned 8-, 16-, 32- and 64-bit)
- ▶ records (structs)
  - ▶ e.g. `{f1 : A, f2 : U32}`
- ▶ variants (tagged unions)
  - ▶ e.g. `<A t1 | B t2>`
- ▶ abstract types
- ▶ function types
- ▶ polymorphic types



Cogent types

# Cogent's linear type system

- ▶ variables with a linear type must be used exactly once
- ▶ COGENT's linear type system:
  - ▶ allows generating efficient imperative code with in-place updates
  - ▶ ensures memory safety

# Limitations

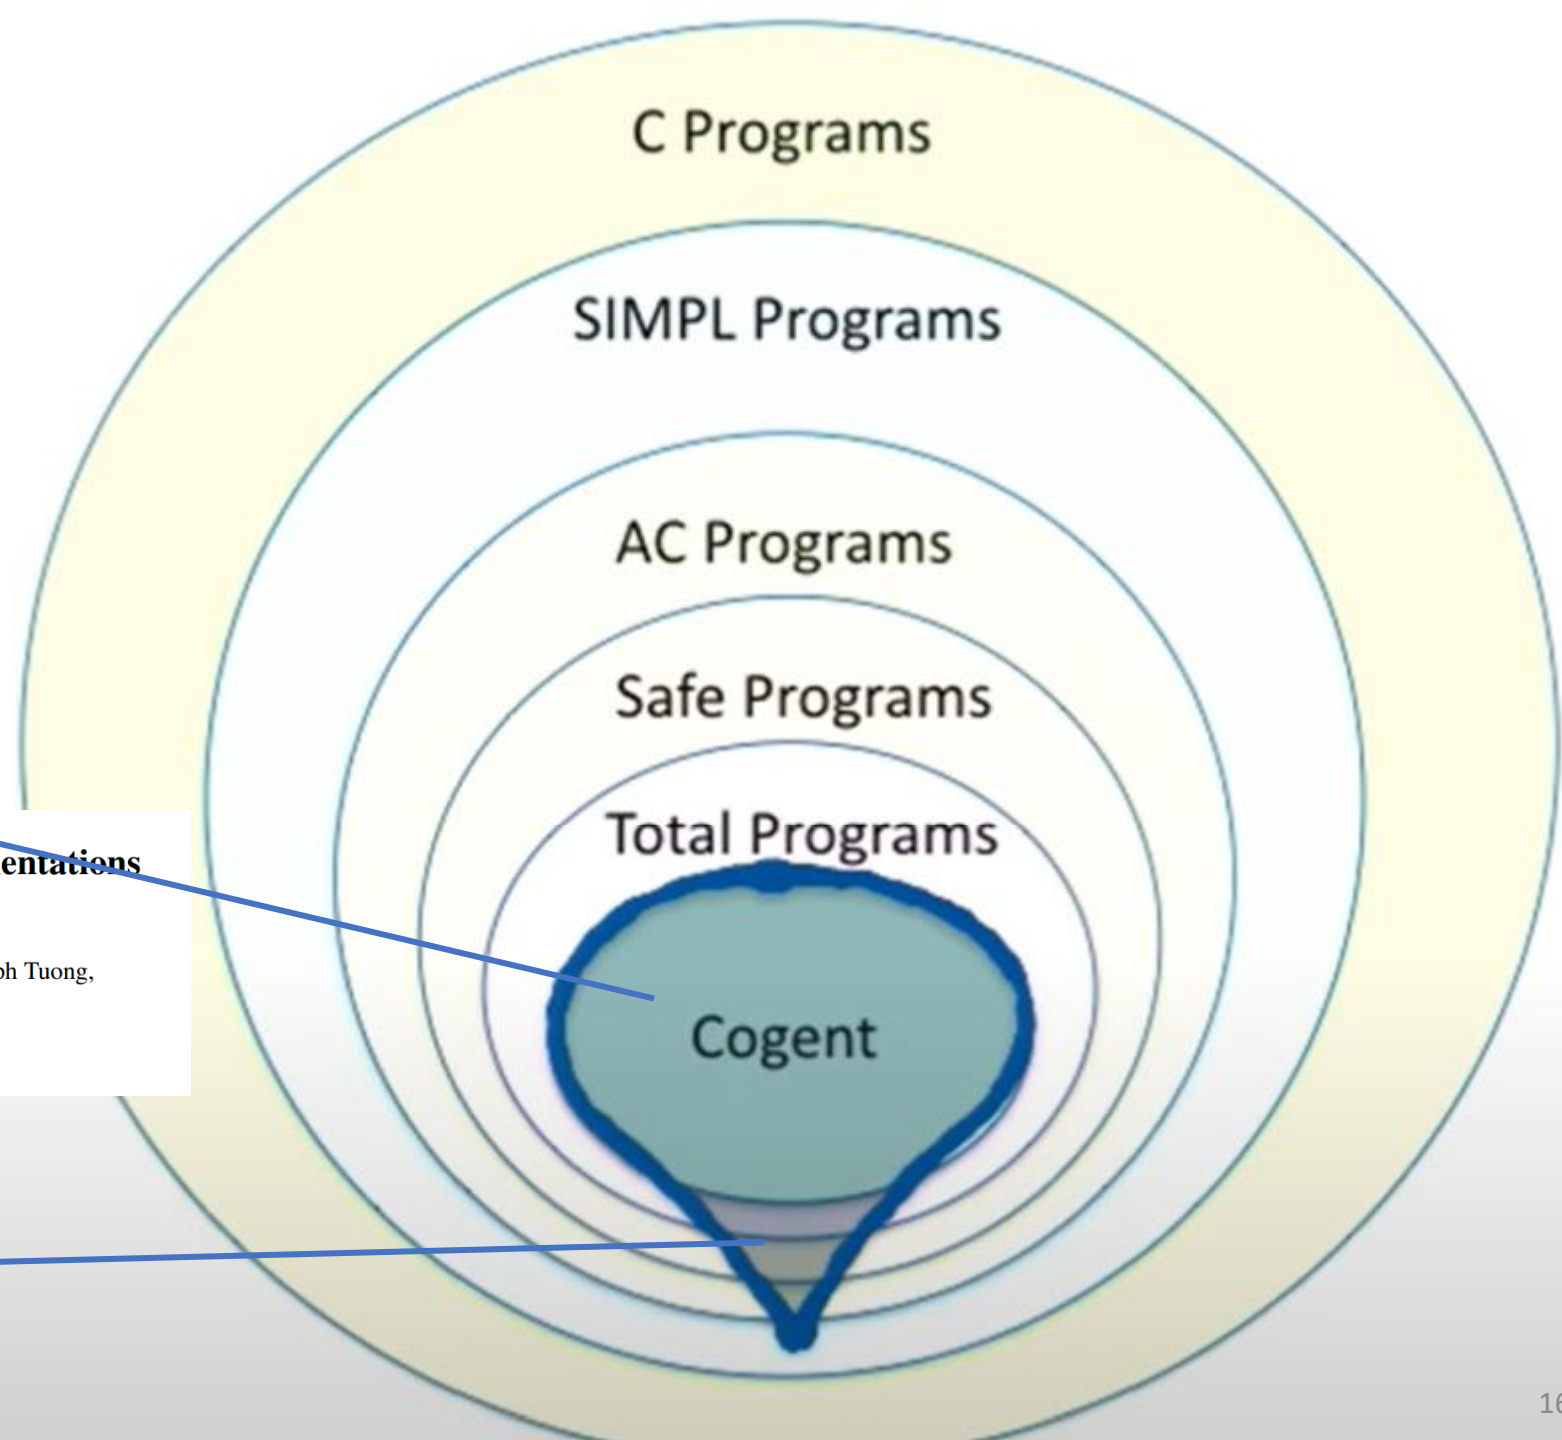
Vast majority of FS-specific code

## COGENT: Verifying High-Assurance File System Implementations

Sidney Amani, Alex Hixon, Zilin Chen, Christine Rizkallah, Peter Chubb,  
Liam O'Connor, Joel Beeren, Yutaka Nagashima, Japheth Lim, Thomas Sewell, Joseph Tuong,  
Gabriele Keller, Toby Murray, Gerwin Klein, Gernot Heiser

Data61 (formerly NICTA) and UNSW, Australia  
[first.last@data61.csiro.au](mailto:first.last@data61.csiro.au)

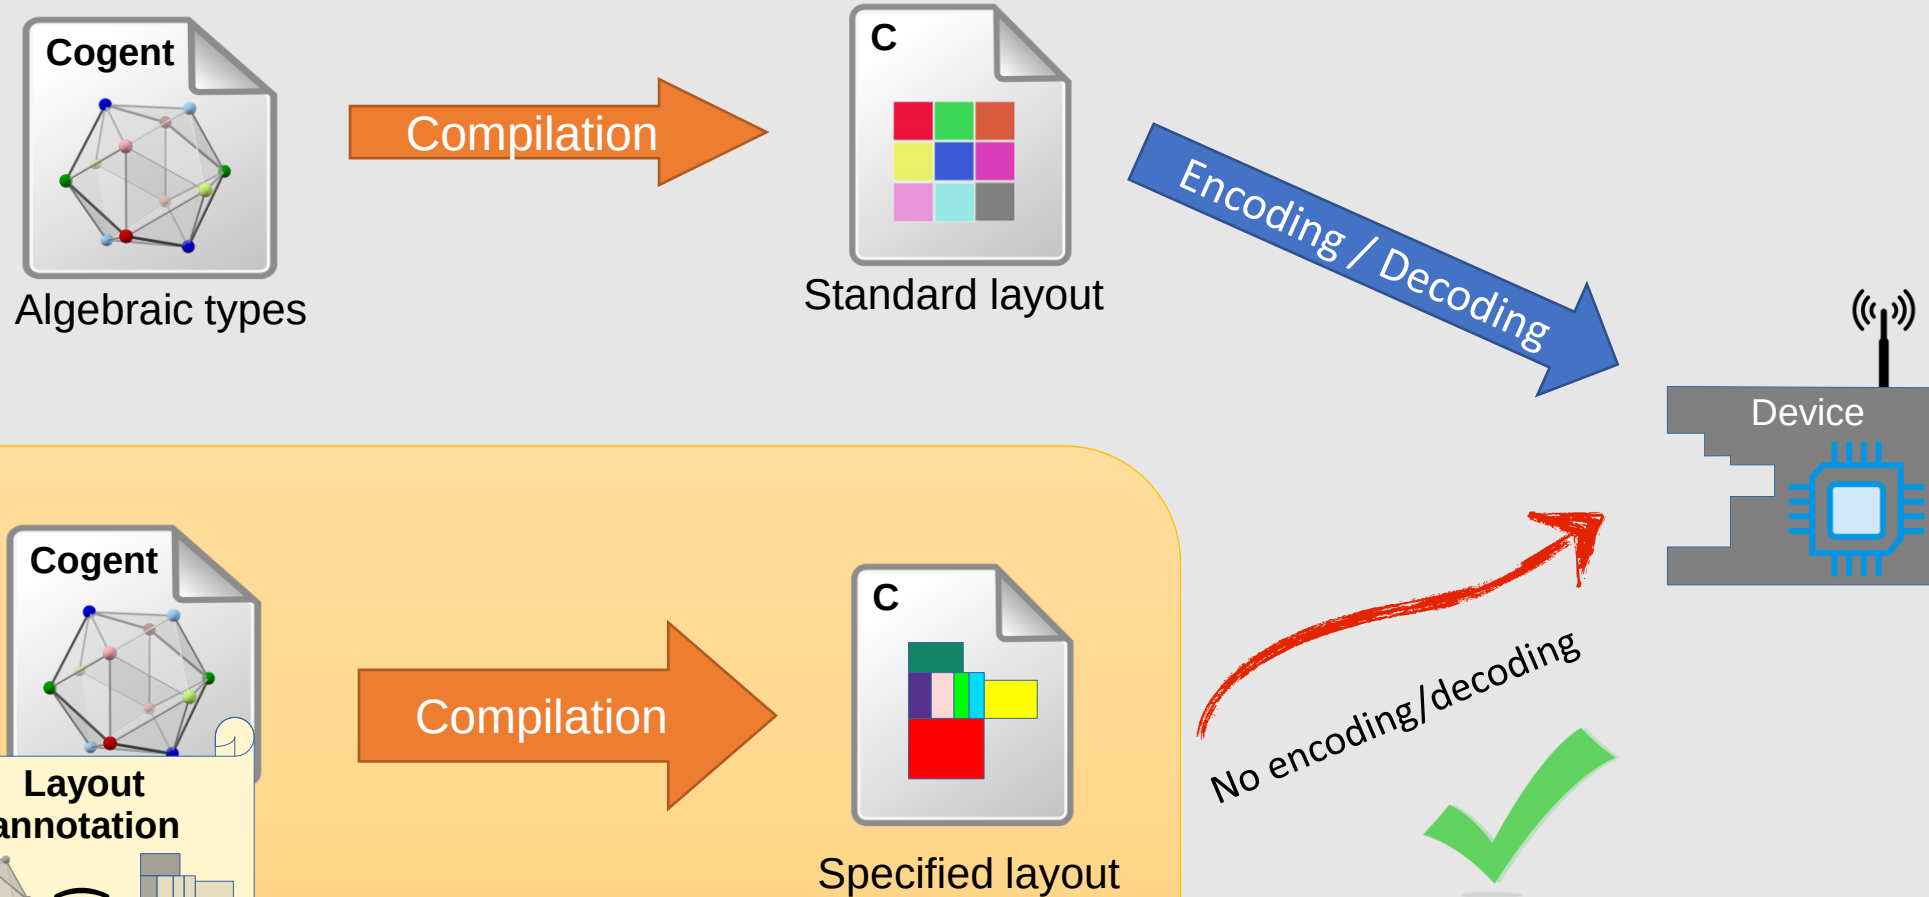
Loop iterators, data  
structures, allocators  
(FFI)





Dargent

# Customising the data layouts



**Dargent Extension**

# Cogent<sup>+</sup> = Cogent + Dargent


Cogent enriched with the possibility to annotate record types with explicit layouts.

Dedicated syntax

Cogent Type

```
type Example = {  
  struct : #{x : U32, y : Bool},  
  ptr : {...},  
  sum : ⟨X U16 | Y U8⟩  
}
```

Possible layout



# Syntax of Dargent annotations

```

type Example = {
  record    struct : #{x : U32, y : Bool},
  pointer   ptr : {...},
  variant   sum : <X U16 | Y U8>
}
    
```

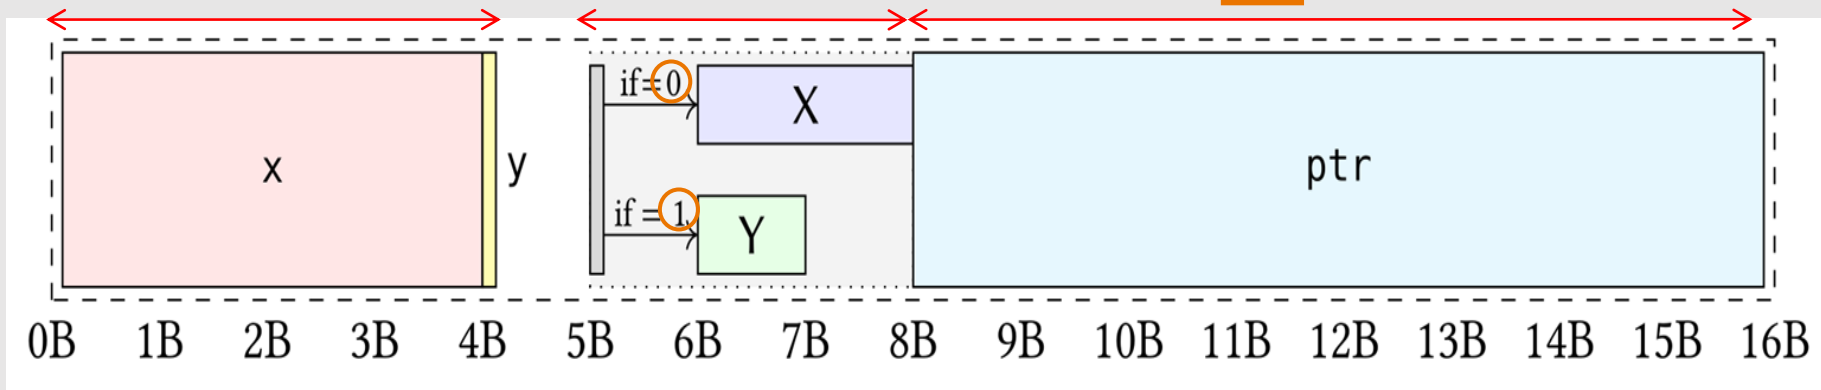
```

layout ExampleLayout = record {
  struct : record { x : 4B, y : 1b },
  ptr : pointer at 8B,
  sum : variant (1b)
        { X(0) : 2B at 1B, Y(1) : 1B at 1B } at 5B
}
    
```

layout size in bits & Bytes

offset

tag



# Simulating Bitfields

## C

```
struct can_id {  
    uint32_t id:29;  
    uint32_t exide:1;  
    uint32_t rtr:1;  
    uint32_t err:1;  
};
```

## Cogent

```
type CanId = {  
    id : U29,  
    exide : Bool,  
    rtr : Bool,  
    err : Bool  
}
```

## Dargent

```
layout record {  
    id : 29b,  
    exide : 1b,  
    rtr : 1b,  
    err : 1b  
}
```

Outside the C parser fragment

# Some features

- Custom endianness (Little / Big endian) for primitive types
- Layout polymorphism

**type** Pair  $t = \{\text{fst} : t, \text{snd} : t\}$

**layout** LPair  $l = \text{record } \{\text{fst} : l, \text{snd} : l \text{ at } 4\text{B}\}$

*freePair* :  $\forall(t, l : \sim t). \text{Pair } t \text{ layout LPair } l \rightarrow ()$

# Matching relation between layouts and types

$$\boxed{\ell \sim \hat{\tau}}$$

$$\frac{\text{bits } (T) = s}{\text{BitRange } (o, s) \sim T} \text{PRIMTYMATCH}$$

$$\begin{aligned} \text{bits } (\text{Un}) &= n \\ \text{bits } (\text{Bool}) &= 1 \end{aligned}$$

$$\frac{M \text{ is the pointer size}}{\text{BitRange } (o, M) \sim \text{pointer}} \text{BOXEDTYMATCH} \qquad \frac{}{x\{o\} \sim \hat{\tau}} \text{VARMATCH}$$

$$\frac{}{() \sim ()} \text{UNITMATCH} \qquad \frac{\text{for each } i: \ell_i \sim \hat{\tau}_i}{\text{record } \{\overline{f_i : \ell_i}\} \sim \{\overline{f_i : \hat{\tau}_i}\}} \text{URECORDMATCH}$$

$$\frac{\text{for each } i: \ell_i \sim \hat{\tau}_i}{\text{variant } (s) \{\overline{A_i (v_i) : \ell_i}\} \sim \langle \overline{A_i \hat{\tau}_i} \rangle} \text{VARIANTMATCH}$$

# Wellformed layouts

$$\boxed{\ell \text{ wf}}$$

$$\frac{\begin{array}{l} \text{for each } i: \ell_i \text{ wf} \\ \text{for each } i \neq j: \text{taken}(\ell_i) \cap \text{taken}(\ell_j) = \emptyset \end{array}}{\text{record } \{\overline{f_i : \ell_i}\} \text{ wf}} \text{URECORDWF}$$

$$\frac{\begin{array}{l} \text{for each } i: \ell_i \text{ wf} \quad v_i < 2^s \quad \text{taken}(\ell_i) \cap \text{taken}(\text{BitRange } (o, s)) = \emptyset \\ \text{for each } i \neq j: v_i \neq v_j \end{array}}{\text{variant } (\text{BitRange } (o, s)) \{\overline{A_i(v_i) : \ell_i}\} \text{ wf}} \text{VARIANTWF}$$

where  $\text{taken}(\ell) \in \mathbb{N}$  returns the set of bit positions that  $\ell$  occupies (defined recursively).



# Compiling Cogent<sup>+</sup> to C

```
type Example = {  
  struct : #{x : U32, y : Bool},  
  ptr : {...},  
  sum : ⟨X U16 | Y U8⟩  
}
```


Compilation

Type : array of 16 bytes

Getters/setters for  
each field

$get_{ptr} : \text{byte}[16] \rightarrow ptr$   
 $set_{ptr} : \text{byte}[16] \times ptr \rightarrow ()$

C





# Example

A power controller system (libopencm3)




`PWR_CR1.VOS = scale`



`set_VOS(PWR_CR1, scale);`

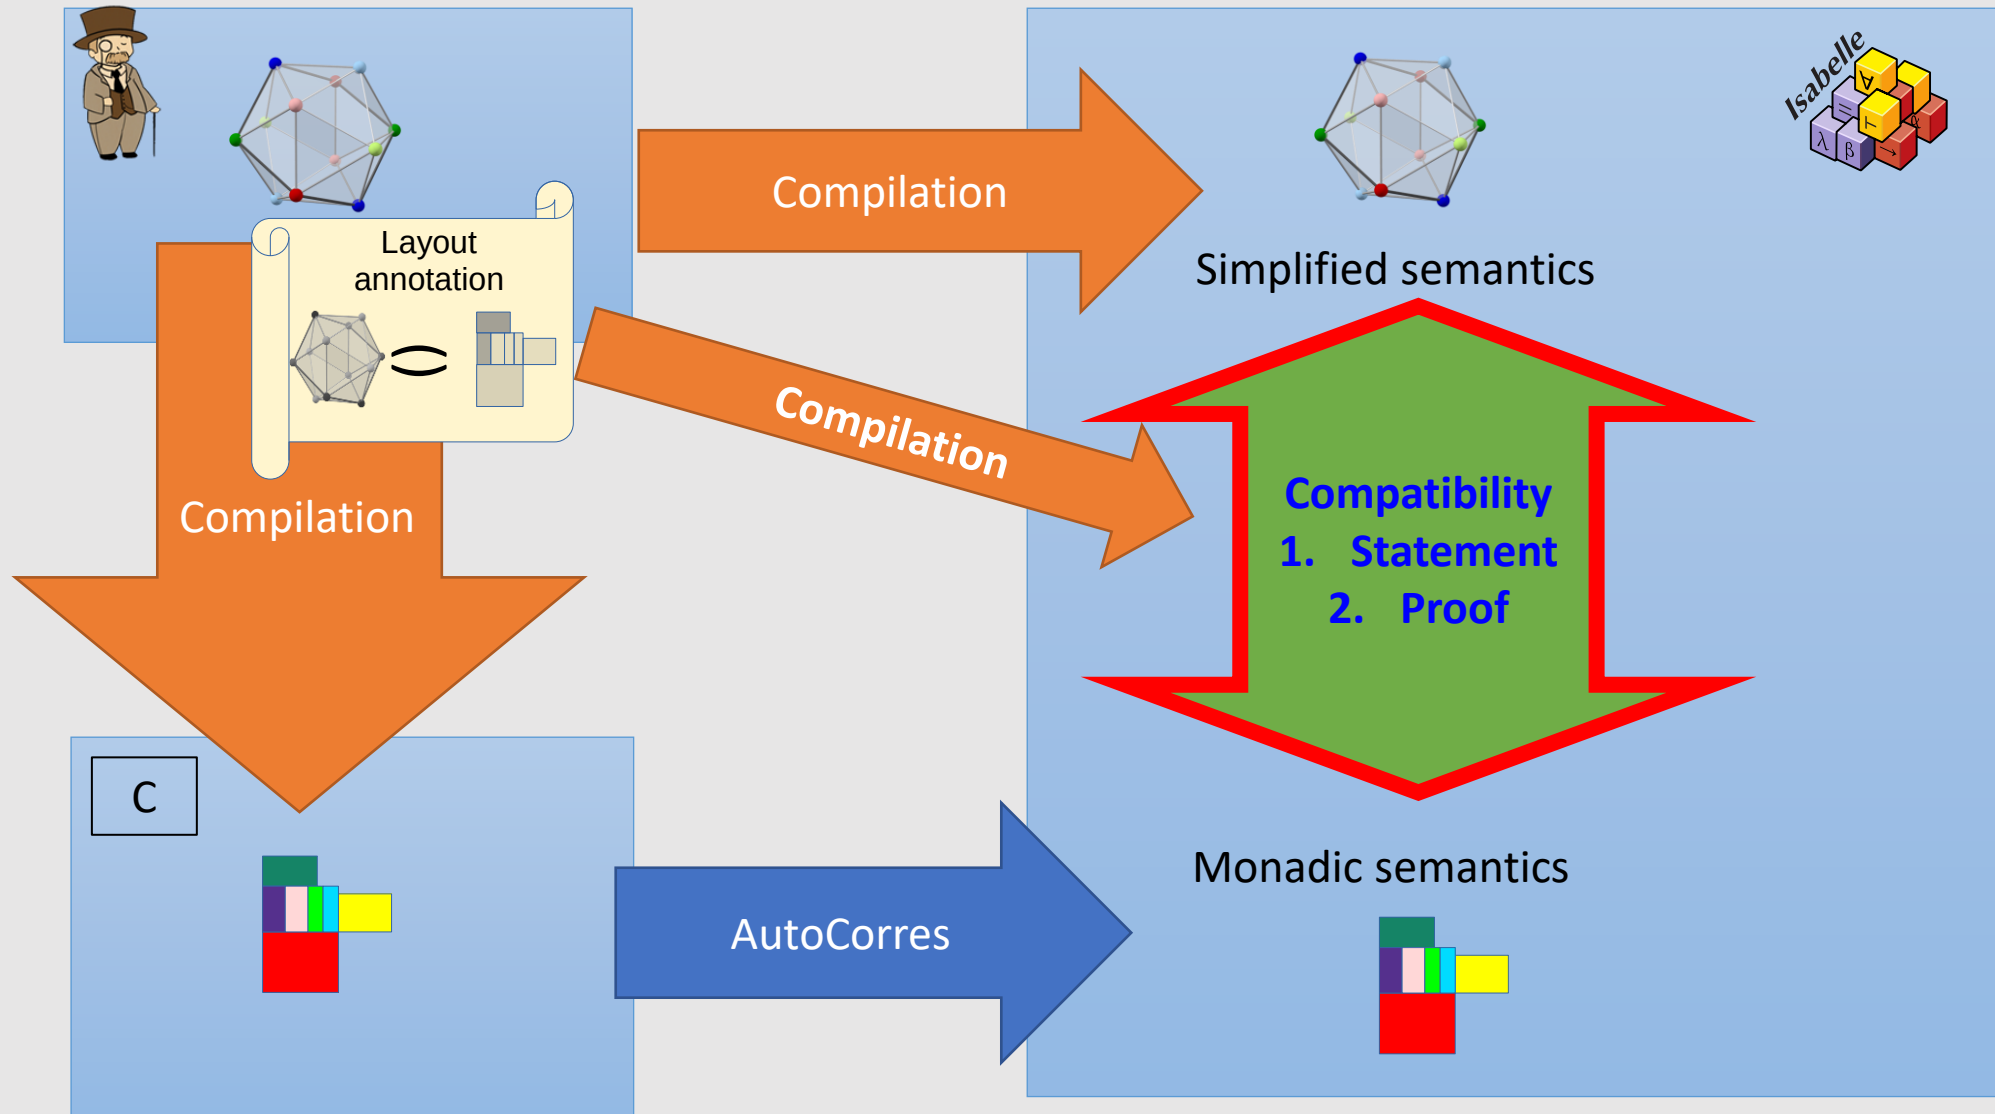
**C**

```
reg32 = PWR_CR1 & ~(PWR_CR1_VOS_MASK << PWR_CR1_VOS_SHIFT);  
reg32 |= (scale & PWR_CR1_VOS_MASK) << PWR_CR1_VOS_SHIFT;  
PWR_CR1 = reg32;
```



Demo

# Compiling Cogent<sup>+</sup>





# Stating compatibility

Monadic semantics	Simplified semantics
Type $T_m$	Type $T_s$
$f_m: A_m \rightarrow B_m$	$f_s: A_s \rightarrow B_s$

Compatibility between the two semantics
Relation $T_r \subset state \times T_m \times T_s$ ?
$(f_m, f_s)$ is compatible with $(A_r, B_r)$

- For an unannotated type, same  $T_r$  as before
- What about an annotated type?

$T_m = \text{byte}[n]$	$T_s = \{ \dots, x : A, \dots \}$
------------------------	-----------------------------------

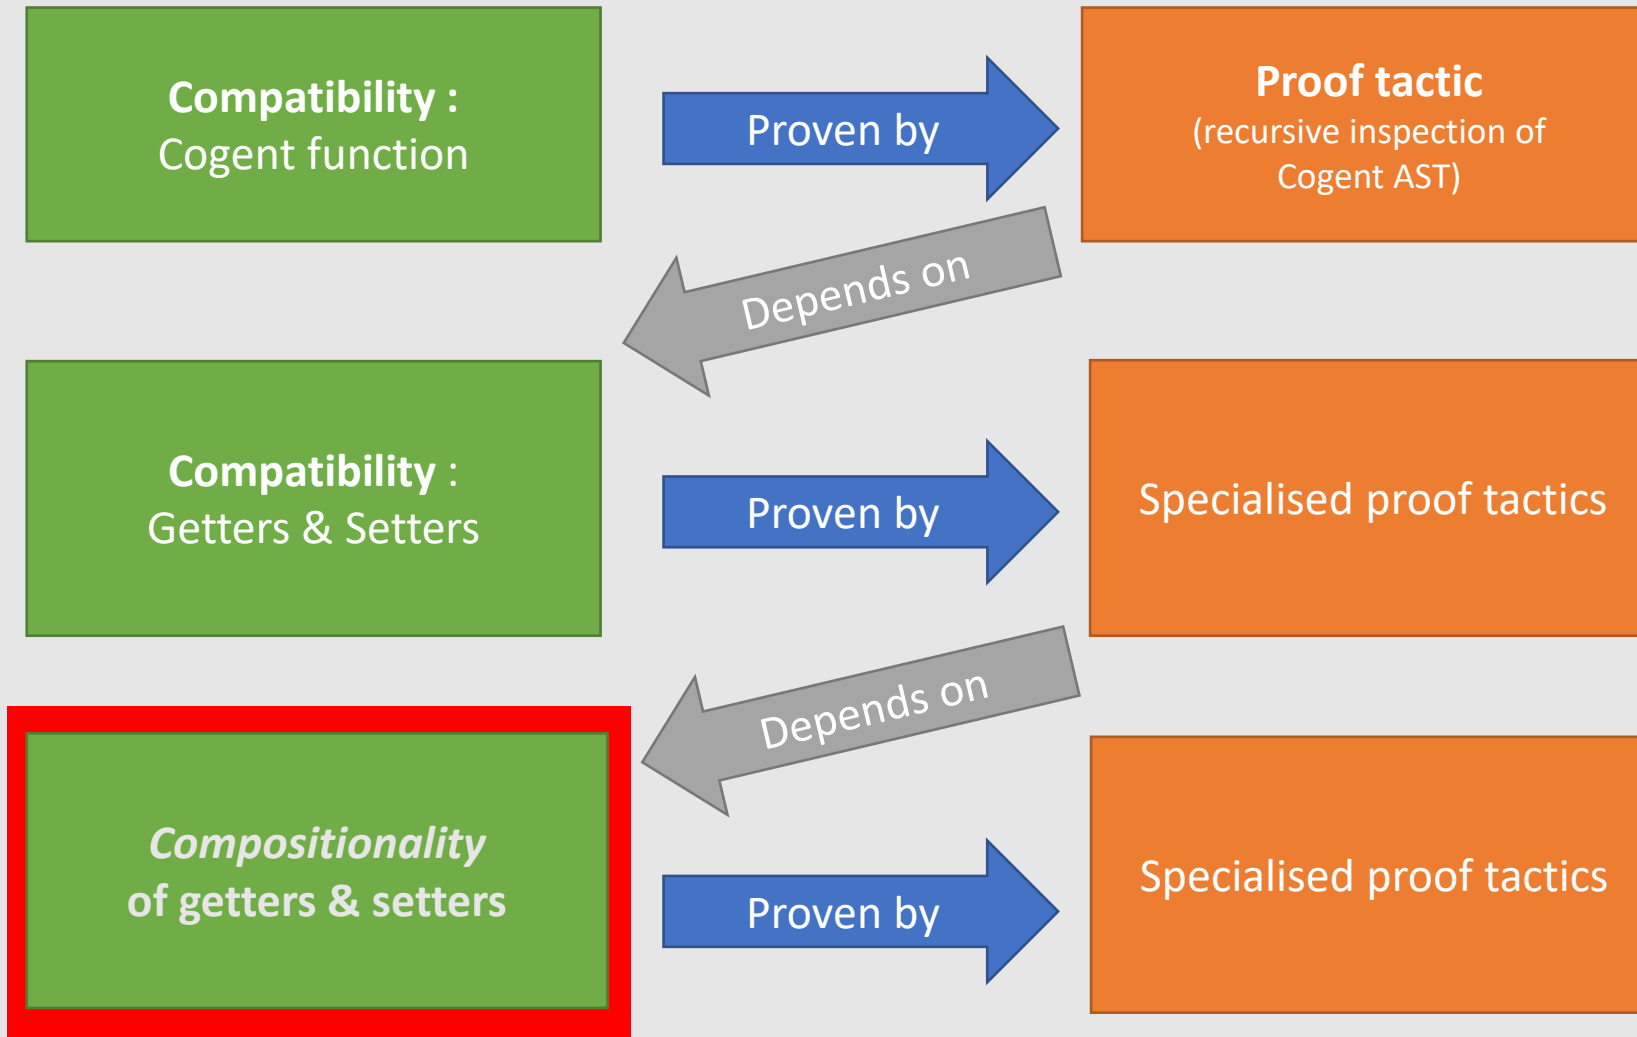


Decoding relation : «  $T_r = \{ (\dots, array, decode(array)) \}$  »

Decodes the array  
according to the layout

$$decode : array \mapsto \left\{ x := \overset{\dots}{get_x}(array) \right\}$$

# Proving compatibility



# Compositionality of getters & setters

$$\ll \begin{array}{l} get_x(set_x(array, value)) = value \\ get_x(set_y(array, value)) = get_x(array) \end{array} \gg$$

Enough to prove compatibility between the two semantics of getters & setters

Monadic semantics	Simplified semantics
$T_m = \text{byte}[n]$	$T_s = \{ \dots, x : A, \dots \}$
$get_x$	$t \mapsto t.x$
$set_x$	$(t, a) \mapsto (t.x := a)$

# Towards an extensible Dargent

Dargent is limited.

For example, what if we want to target something else than a byte array?

## **Idea :**

Instead of describing the layout, the user provides

- The target type (e.g., a byte array)
- An implementation of getters / setters
- The proofs of their compositionality properties

Enough to ensure compatibility!

**Demo**

# Certifying getters & setters

## **Observation :**

Compatibility between the semantics gives no guarantee that the generated getters / setters conform to the specified layout!

## **Additional generated formal guarantees**

1. Generated setters do not flip any bit outside their layouts
2. Generated getters specialise a generic getter implemented in Isabelle (parameterised by a layout)

**Demo**



# Some feature wishes

- Recursive types, Primitive arrays (Cogent)
- Controlling the layout of the field type in getters / setters

$$\begin{array}{ccc} T = \{ \dots, x : A, \dots \} & \begin{array}{l} get_x: byte[n] \rightarrow A \\ set_x: byte[n] \times A \rightarrow () \end{array} & \text{VS} \quad \begin{array}{l} get_x: byte[n] \rightarrow byte[n_A] \\ set_x: byte[n] \times byte[n_A] \rightarrow () \end{array} \end{array}$$

**Example :**  $A = \langle X \ U8 \mid Y \ U16 \rangle$

$$\begin{array}{ccc} get_x: byte[n] \rightarrow \left\{ \begin{array}{l} tag: U8 \\ payload_x: U8 \\ payload_y: U16 \end{array} \right\} & \text{VS} & get_x: byte[n] \rightarrow byte[3] \end{array}$$

# Some feature wishes

- Simultaneous field updates

```
regs {  
    timer_a_en = True,  
    timer_a_input_clk =  
        TIMEOUT_TIMEBASE_1_MS,  
    timer_e_input_clk =  
        TIMESTAMP_TIMEBASE_1_US }
```



Compiles to?

```
timer->regs->mux =  
    TIMER_A_EN  
    | (TIMEOUT_TIMEBASE_1_MS <<  
        TIMER_A_INPUT_CLK)  
    | (TIMESTAMP_TIMEBASE_1_US <<  
        TIMER_E_INPUT_CLK) ;
```

# A formally verified timer driver

**Demo**

# Conclusion

Cogent + Dargent = First programming language with

1. Possible customisation of the layout of compiled types
2. Formal guarantees

Check our POPL '23 article:

**Dargent: A Silver Bullet for Verified Data Layout Refinement**